




Multilayer Networks


Natural Language Processing: Jordan
Boyd-Graber
University of Maryland

SLIDES ADAPTED FROM ANDREW NG

Logistic Regression by Another Name: Map inputs to output



Logistic Regression by Another Name: Map inputs to output




Input

Vector $x_1 \dots x_d$

inputs encoded as
real numbers

Logistic Regression by Another Name: Map inputs to output



Output


$$f\left(\sum_i w_i x_i + b\right)$$

Input

Vector $x_1 \dots x_d$

multiply inputs by

Logistic Regression by Another Name: Map inputs to output



Input


Vector $x_1 \dots x_d$

Output

$$f\left(\sum_i w_i x_i + b\right)$$

add bias

Logistic Regression by Another Name: Map inputs to output



Input

Vector $x_1 \dots x_d$

Output


$$f\left(\sum_i w_i x_i + b\right)$$

Activation

$$f(z) \equiv \frac{1}{1 + \exp(-z)}$$

pass through
nonlinear sigmoid


Why is it called activation?



In the shallow end

- This is still logistic regression
- Engineering features x is difficult (and requires expertise)
- Can we learn how to represent inputs into final decision?

Better name: non-linearity



- Logistic / Sigmoid

$$f(x) = \frac{1}{1 + e^{-x}} \quad (1)$$

- tanh

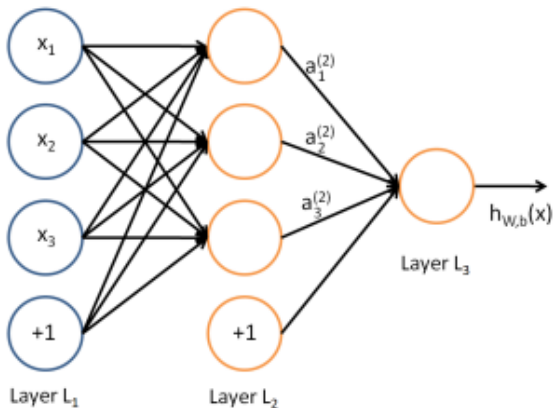
$$f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1 \quad (2)$$

- ReLU

$$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases} \quad (3)$$


- SoftPlus: $f(x) = \ln(1 + e^x)$

Learn the features and the function




$$a_1^{(2)} = f\left(W_{11}^{(1)} x_1 + W_{12}^{(1)} x_2 + W_{13}^{(1)} x_3 + b_1^{(1)}\right)$$

Learn the features and the function



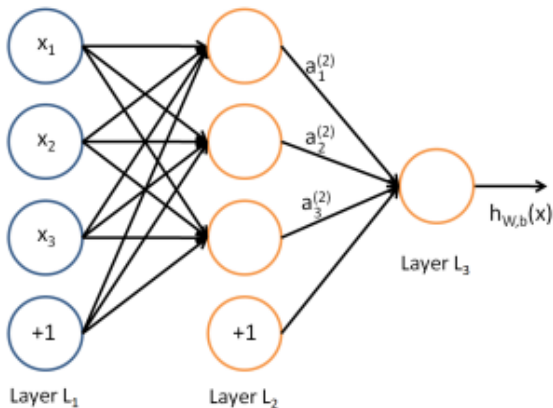
$$a_2^{(2)} = f\left(W_{21}^{(1)}x_1 + W_{22}^{(1)}x_2 + W_{23}^{(1)}x_3 + b_2^{(1)}\right)$$

Learn the features and the function



$$a_3^{(2)} = f\left(w_{31}^{(1)}x_1 + w_{32}^{(1)}x_2 + w_{33}^{(1)}x_3 + b_3^{(1)}\right)$$

Learn the features and the function



$$h_{W,b}(x) = a_1^{(3)} = f\left(W_{11}^{(2)} a_1^{(2)} + W_{12}^{(2)} a_2^{(2)} + W_{13}^{(2)} a_3^{(2)} + b_1^{(2)}\right)$$

Objective Function

- For every example x, y of our supervised training set, we want the label y to match the prediction $h_{W,b}(x)$.

$$J(W, b; x, y) \equiv \frac{1}{2} \|h_{W,b}(x) - y\|^2 \quad (4)$$

Objective Function

- For every example x, y of our supervised training set, we want the label y to match the prediction $h_{W,b}(x)$.

$$J(W, b; x, y) \equiv \frac{1}{2} \|h_{W,b}(x) - y\|^2 \quad (4)$$

- We want this value, summed over all of the examples to be as small as possible

Objective Function

- For every example x, y of our supervised training set, we want the label y to match the prediction $h_{W,b}(x)$.

$$J(W, b; x, y) \equiv \frac{1}{2} \|h_{W,b}(x) - y\|^2 \quad (4)$$

- We want this value, summed over all of the examples to be as small as possible
- We also want the weights not to be too large

$$\frac{\lambda}{2} \sum_l^{n_l-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (w_{ji}^l)^2 \quad (5)$$

Objective Function

- For every example x, y of our supervised training set, we want the label y to match the prediction $h_{W,b}(x)$.

$$J(W, b; x, y) \equiv \frac{1}{2} \|h_{W,b}(x) - y\|^2 \quad (4)$$

- We want this value, summed over all of the examples to be as small as possible
- We also want the weights not to be too large

$$\frac{\lambda}{2} \sum_l^{n_l-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (w_{ji}^l)^2 \quad (5)$$

Objective Function

- For every example x, y of our supervised training set, we want the label y to match the prediction $h_{W,b}(x)$.

$$J(W, b; x, y) \equiv \frac{1}{2} \|h_{W,b}(x) - y\|^2 \quad (4)$$

- We want this value, summed over all of the examples to be as small as possible
- We also want the weights not to be too large

$$\frac{\lambda}{2} \sum_l^{n_l-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (w_{ji}^l)^2 \quad (5)$$

Sum over all layers

Objective Function

- For every example x, y of our supervised training set, we want the label y to match the prediction $h_{W,b}(x)$.

$$J(W, b; x, y) \equiv \frac{1}{2} \|h_{W,b}(x) - y\|^2 \quad (4)$$

- We want this value, summed over all of the examples to be as small as possible
- We also want the weights not to be too large

$$\frac{\lambda}{2} \sum_l^{n_l-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (w_{ji}^l)^2 \quad (5)$$

Sum over all sources

Objective Function

- For every example x, y of our supervised training set, we want the label y to match the prediction $h_{W,b}(x)$.

$$J(W, b; x, y) \equiv \frac{1}{2} \|h_{W,b}(x) - y\|^2 \quad (4)$$

- We want this value, summed over all of the examples to be as small as possible
- We also want the weights not to be too large

$$\frac{\lambda}{2} \sum_l^{n_l-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (w_{ji}^l)^2 \quad (5)$$

Sum over all destinations

Objective Function

Putting it all together:

$$J(W, b) = \left[\frac{1}{m} \sum_{i=1}^m \frac{1}{2} \|h_{W,b}(x^{(i)}) - y^{(i)}\|^2 \right] + \frac{\lambda}{2} \sum_l^{n_l-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (w_{ji}^l)^2 \quad (6)$$

Objective Function

Putting it all together:

$$J(W, b) = \left[\frac{1}{m} \sum_{i=1}^m \frac{1}{2} \|h_{W,b}(x^{(i)}) - y^{(i)}\|^2 \right] + \frac{\lambda}{2} \sum_l^{n_l-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (W_{ji}^l)^2 \quad (6)$$

- Our goal is to minimize $J(W, b)$ as a function of W and b

Objective Function

Putting it all together:

$$J(W, b) = \left[\frac{1}{m} \sum_{i=1}^m \frac{1}{2} \|h_{W,b}(x^{(i)}) - y^{(i)}\|^2 \right] + \frac{\lambda}{2} \sum_l^{n_l-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (w_{ji}^l)^2 \quad (6)$$

- Our goal is to minimize $J(W, b)$ as a function of W and b
- Initialize W and b to small random value near zero

Objective Function

Putting it all together:


$$J(W, b) = \left[\frac{1}{m} \sum_{i=1}^m \frac{1}{2} \|h_{W,b}(x^{(i)}) - y^{(i)}\|^2 \right] + \frac{\lambda}{2} \sum_l^{n_l-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (w_{ji}^l)^2 \quad (6)$$

- Our goal is to minimize $J(W, b)$ as a function of W and b
- Initialize W and b to small random value near zero
- Adjust parameters to optimize J

Gradient Descent

Goal

Optimize J with respect to variables W and b



Backpropigation

- For convenience, write the input to sigmoid

$$z_i^{(l)} = \sum_{j=1}^n w_{ij}^{(l-1)} x_j + b_i^{(l-1)} \quad (7)$$

Backpropigation

- For convenience, write the input to sigmoid

$$z_i^{(l)} = \sum_{j=1}^n w_{ij}^{(l-1)} x_j + b_i^{(l-1)} \quad (7)$$

- The gradient is a function of a node's error $\delta_i^{(l)}$

Backpropigation

- For convenience, write the input to sigmoid

$$z_i^{(l)} = \sum_{j=1}^n w_{ij}^{(l-1)} x_j + b_i^{(l-1)} \quad (7)$$

- The gradient is a function of a node's error $\delta_i^{(l)}$
- For output nodes, the error is obvious:

$$\delta_i^{(n_l)} = \frac{\partial}{\partial z_i^{(n_l)}} \|y - h_{w,b}(x)\|^2 = -(y_i - a_i^{(n_l)}) \cdot f'(z_i^{(n_l)}) \frac{2}{2} \quad (8)$$

Backpropagation

- For convenience, write the input to sigmoid

$$z_i^{(l)} = \sum_{j=1}^n w_{ij}^{(l-1)} x_j + b_i^{(l-1)} \quad (7)$$

- The gradient is a function of a node's error $\delta_i^{(l)}$
- For output nodes, the error is obvious:

$$\delta_i^{(n_l)} = \frac{\partial}{\partial z_i^{(n_l)}} \|y - h_{w,b}(x)\|^2 = -(y_i - a_i^{(n_l)}) \cdot f'(z_i^{(n_l)}) \frac{2}{2} \quad (8)$$

- Other nodes must “backpropagate” **downstream error** based on connection strength

$$\delta_i^{(l)} = \left(\sum_{j=1}^{s_{l+1}} w_{ji}^{(l+1)} \delta_j^{(l+1)} \right) f'(z_i^{(l)}) \quad (9)$$

Backpropagation

- For convenience, write the input to sigmoid

$$z_i^{(l)} = \sum_{j=1}^n w_{ij}^{(l-1)} x_j + b_i^{(l-1)} \quad (7)$$

- The gradient is a function of a node's error $\delta_i^{(l)}$
- For output nodes, the error is obvious:

$$\delta_i^{(n_l)} = \frac{\partial}{\partial z_i^{(n_l)}} \|y - h_{w,b}(x)\|^2 = -(y_i - a_i^{(n_l)}) \cdot f'(z_i^{(n_l)}) \frac{2}{2} \quad (8)$$

- Other nodes must “backpropagate” downstream error based on **connection strength**

$$\delta_i^{(l)} = \left(\sum_{j=1}^{s_{l+1}} w_{ji}^{(l+1)} \delta_j^{(l+1)} \right) f'(z_i^{(l)}) \quad (9)$$

Backpropagation

- For convenience, write the input to sigmoid

$$z_i^{(l)} = \sum_{j=1}^n w_{ij}^{(l-1)} x_j + b_i^{(l-1)} \quad (7)$$

- The gradient is a function of a node's error $\delta_i^{(l)}$
- For output nodes, the error is obvious:

$$\delta_i^{(n_l)} = \frac{\partial}{\partial z_i^{(n_l)}} \|y - h_{w,b}(x)\|^2 = -(y_i - a_i^{(n_l)}) \cdot f'(z_i^{(n_l)}) \frac{2}{2} \quad (8)$$

- Other nodes must “backpropagate” downstream error based on connection strength

$$\delta_i^{(l)} = \left(\sum_{j=1}^{s_{l+1}} w_{ji}^{(l+1)} \delta_j^{(l+1)} \right) f'(z_i^{(l)}) \quad (9)$$

(chain rule)

Partial Derivatives

- For weights, the partial derivatives are

$$\frac{\partial}{\partial W_{ij}^{(l)}} J(W, b; x, y) = a_j^{(l)} \delta_i^{(l+1)} \quad (10)$$

- For the bias terms, the partial derivatives are

$$\frac{\partial}{\partial b_i^{(l)}} J(W, b; x, y) = \delta_i^{(l+1)} \quad (11)$$

- But this is just for a single example ...

Full Gradient Descent Algorithm

1. Initialize $U^{(l)}$ and $V^{(l)}$ as zero
2. For each example $i = 1 \dots m$
 - 2.1 Use backpropagation to compute $\nabla_W J$ and $\nabla_b J$
 - 2.2 Update weight shifts $U^{(l)} = U^{(l)} + \nabla_{W^{(l)}} J(W, b; x, y)$
 - 2.3 Update bias shifts $V^{(l)} = V^{(l)} + \nabla_{b^{(l)}} J(W, b; x, y)$
3. Update the parameters

$$W^{(l)} = W^{(l)} - \alpha \left[\left(\frac{1}{m} U^{(l)} \right) \right] \quad (12)$$

$$b^{(l)} = b^{(l)} - \alpha \left[\frac{1}{m} V^{(l)} \right] \quad (13)$$

4. Repeat until weights stop changing

But it is not perfect

- Compare against baselines: randomized features, nearest-neighbors, linear models
- Optimization is hard (alchemy)
- Models are often not interpretable
- Requires specialized hardware and tons of data to scale




Multilayer Networks

Natural Language Processing: Jordan
Boyd-Graber

University of Maryland


MATHEMATICAL DESCRIPTION

Learn the features and the function




$$a_1^{(2)} = f\left(W_{11}^{(1)}x_1 + W_{12}^{(1)}x_2 + W_{13}^{(1)}x_3 + b_1^{(1)}\right)$$

Learn the features and the function




$$a_2^{(2)} = f\left(W_{21}^{(1)}x_1 + W_{22}^{(1)}x_2 + W_{23}^{(1)}x_3 + b_2^{(1)}\right)$$

Learn the features and the function



$$a_3^{(2)} = f\left(w_{31}^{(1)}x_1 + w_{32}^{(1)}x_2 + w_{33}^{(1)}x_3 + b_3^{(1)}\right)$$

Learn the features and the function



$$h_{W,b}(x) = a_1^{(3)} = f\left(W_{11}^{(2)} a_1^{(2)} + W_{12}^{(2)} a_2^{(2)} + W_{13}^{(2)} a_3^{(2)} + b_1^{(2)}\right)$$

Objective Function

- For every example x, y of our supervised training set, we want the label y to match the prediction $h_{W,b}(x)$.

$$J(W, b; x, y) \equiv \frac{1}{2} \|h_{W,b}(x) - y\|^2 \quad (1)$$

Objective Function

- For every example x, y of our supervised training set, we want the label y to match the prediction $h_{W,b}(x)$.

$$J(W, b; x, y) \equiv \frac{1}{2} \|h_{W,b}(x) - y\|^2 \quad (1)$$

- We want this value, summed over all of the examples to be as small as possible

Objective Function

- For every example x, y of our supervised training set, we want the label y to match the prediction $h_{W,b}(x)$.

$$J(W, b; x, y) \equiv \frac{1}{2} \|h_{W,b}(x) - y\|^2 \quad (1)$$

- We want this value, summed over all of the examples to be as small as possible
- We also want the weights not to be too large

$$\frac{\lambda}{2} \sum_l^{n_l-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (w_{ji}^l)^2 \quad (2)$$

Objective Function

- For every example x, y of our supervised training set, we want the label y to match the prediction $h_{W,b}(x)$.

$$J(W, b; x, y) \equiv \frac{1}{2} \|h_{W,b}(x) - y\|^2 \quad (1)$$

- We want this value, summed over all of the examples to be as small as possible
- We also want the weights not to be too large

$$\frac{\lambda}{2} \sum_l^{n_l-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (w_{ji}^l)^2 \quad (2)$$

Objective Function

- For every example x, y of our supervised training set, we want the label y to match the prediction $h_{W,b}(x)$.

$$J(W, b; x, y) \equiv \frac{1}{2} \|h_{W,b}(x) - y\|^2 \quad (1)$$

- We want this value, summed over all of the examples to be as small as possible
- We also want the weights not to be too large

$$\frac{\lambda}{2} \sum_l^{n_l-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (w_{ji}^l)^2 \quad (2)$$

Sum over all layers

Objective Function

- For every example x, y of our supervised training set, we want the label y to match the prediction $h_{W,b}(x)$.

$$J(W, b; x, y) \equiv \frac{1}{2} \|h_{W,b}(x) - y\|^2 \quad (1)$$

- We want this value, summed over all of the examples to be as small as possible
- We also want the weights not to be too large

$$\frac{\lambda}{2} \sum_l^{n_l-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (w_{ji}^l)^2 \quad (2)$$

Sum over all sources

Objective Function

- For every example x, y of our supervised training set, we want the label y to match the prediction $h_{W,b}(x)$.

$$J(W, b; x, y) \equiv \frac{1}{2} \|h_{W,b}(x) - y\|^2 \quad (1)$$

- We want this value, summed over all of the examples to be as small as possible
- We also want the weights not to be too large

$$\frac{\lambda}{2} \sum_l^{n_l-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (w_{ji}^l)^2 \quad (2)$$

Sum over all destinations

Objective Function

Putting it all together:

$$J(W, b) = \left[\frac{1}{m} \sum_{i=1}^m \frac{1}{2} \|h_{W,b}(x^{(i)}) - y^{(i)}\|^2 \right] + \frac{\lambda}{2} \sum_l^{n_l-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (w_{ji}^l)^2 \quad (3)$$

Objective Function

Putting it all together:

$$J(W, b) = \left[\frac{1}{m} \sum_{i=1}^m \frac{1}{2} \|h_{W,b}(x^{(i)}) - y^{(i)}\|^2 \right] + \frac{\lambda}{2} \sum_l^{n_l-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (w_{ji}^l)^2 \quad (3)$$

- Our goal is to minimize $J(W, b)$ as a function of W and b

Objective Function

Putting it all together:

$$J(W, b) = \left[\frac{1}{m} \sum_{i=1}^m \frac{1}{2} \|h_{W,b}(x^{(i)}) - y^{(i)}\|^2 \right] + \frac{\lambda}{2} \sum_l^{n_l-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (w_{ji}^l)^2 \quad (3)$$

- Our goal is to minimize $J(W, b)$ as a function of W and b
- Initialize W and b to small random value near zero

Objective Function

Putting it all together:


$$J(W, b) = \left[\frac{1}{m} \sum_{i=1}^m \frac{1}{2} \|h_{W,b}(x^{(i)}) - y^{(i)}\|^2 \right] + \frac{\lambda}{2} \sum_l^{n_l-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (w_{ji}^l)^2 \quad (3)$$

- Our goal is to minimize $J(W, b)$ as a function of W and b
- Initialize W and b to small random value near zero
- Adjust parameters to optimize J

Gradient Descent

Goal

Optimize J with respect to variables W and b



Backpropigation

- For convenience, write the input to sigmoid

$$z_i^{(l)} = \sum_{j=1}^n w_{ij}^{(l-1)} x_j + b_i^{(l-1)} \quad (4)$$

Backpropigation

- For convenience, write the input to sigmoid

$$z_i^{(l)} = \sum_{j=1}^n w_{ij}^{(l-1)} x_j + b_i^{(l-1)} \quad (4)$$

- The gradient is a function of a node's error $\delta_i^{(l)}$

Backpropigation

- For convenience, write the input to sigmoid

$$z_i^{(l)} = \sum_{j=1}^n w_{ij}^{(l-1)} x_j + b_i^{(l-1)} \quad (4)$$

- The gradient is a function of a node's error $\delta_i^{(l)}$
- For output nodes, the error is obvious:

$$\delta_i^{(n_l)} = \frac{\partial}{\partial z_i^{(n_l)}} \|y - h_{w,b}(x)\|^2 = -(y_i - a_i^{(n_l)}) \cdot f'(z_i^{(n_l)}) \frac{1}{2} \quad (5)$$

Backpropagation

- For convenience, write the input to sigmoid

$$z_i^{(l)} = \sum_{j=1}^n w_{ij}^{(l-1)} x_j + b_i^{(l-1)} \quad (4)$$

- The gradient is a function of a node's error $\delta_i^{(l)}$
- For output nodes, the error is obvious:

$$\delta_i^{(n_l)} = \frac{\partial}{\partial z_i^{(n_l)}} \|y - h_{w,b}(x)\|^2 = -(y_i - a_i^{(n_l)}) \cdot f'(z_i^{(n_l)}) \frac{1}{2} \quad (5)$$

- Other nodes must “backpropagate” **downstream error** based on connection strength

$$\delta_i^{(l)} = \left(\sum_{j=1}^{s_{l+1}} w_{ji}^{(l+1)} \delta_j^{(l+1)} \right) f'(z_i^{(l)}) \quad (6)$$

Backpropagation

- For convenience, write the input to sigmoid

$$z_i^{(l)} = \sum_{j=1}^n w_{ij}^{(l-1)} x_j + b_i^{(l-1)} \quad (4)$$

- The gradient is a function of a node's error $\delta_i^{(l)}$
- For output nodes, the error is obvious:

$$\delta_i^{(n_l)} = \frac{\partial}{\partial z_i^{(n_l)}} \|y - h_{w,b}(x)\|^2 = -(y_i - a_i^{(n_l)}) \cdot f'(z_i^{(n_l)}) \frac{1}{2} \quad (5)$$

- Other nodes must “backpropagate” downstream error based on **connection strength**

$$\delta_i^{(l)} = \left(\sum_{j=1}^{s_{l+1}} w_{ji}^{(l+1)} \delta_j^{(l+1)} \right) f'(z_i^{(l)}) \quad (6)$$

Backpropagation

- For convenience, write the input to sigmoid

$$z_i^{(l)} = \sum_{j=1}^n w_{ij}^{(l-1)} x_j + b_i^{(l-1)} \quad (4)$$

- The gradient is a function of a node's error $\delta_i^{(l)}$
- For output nodes, the error is obvious:

$$\delta_i^{(n_l)} = \frac{\partial}{\partial z_i^{(n_l)}} \|y - h_{w,b}(x)\|^2 = -(y_i - a_i^{(n_l)}) \cdot f'(z_i^{(n_l)}) \frac{1}{2} \quad (5)$$

- Other nodes must “backpropagate” downstream error based on connection strength

$$\delta_i^{(l)} = \left(\sum_{j=1}^{s_{l+1}} w_{ji}^{(l+1)} \delta_j^{(l+1)} \right) f'(z_i^{(l)}) \quad (6)$$

(chain rule)

Partial Derivatives

- For weights, the partial derivatives are

$$\frac{\partial}{\partial W_{ij}^{(l)}} J(W, b; x, y) = a_j^{(l)} \delta_i^{(l+1)} \quad (7)$$

- For the bias terms, the partial derivatives are

$$\frac{\partial}{\partial b_i^{(l)}} J(W, b; x, y) = \delta_i^{(l+1)} \quad (8)$$

- But this is just for a single example ...

Full Gradient Descent Algorithm

1. Initialize $U^{(l)}$ and $V^{(l)}$ as zero
2. For each example $i = 1 \dots m$
 - 2.1 Use backpropagation to compute $\nabla_W J$ and $\nabla_b J$
 - 2.2 Update weight shifts $U^{(l)} = U^{(l)} + \nabla_{W^{(l)}} J(W, b; x, y)$
 - 2.3 Update bias shifts $V^{(l)} = V^{(l)} + \nabla_{b^{(l)}} J(W, b; x, y)$
3. Update the parameters

$$W^{(l)} = W^{(l)} - \alpha \left[\left(\frac{1}{m} U^{(l)} \right) \right] \quad (9)$$

$$b^{(l)} = b^{(l)} - \alpha \left[\frac{1}{m} V^{(l)} \right] \quad (10)$$

4. Repeat until weights stop changing



Frameworks

Natural Language Processing: Jordan
Boyd-Graber

University of Maryland

INTRODUCTION

Slides adapted from Chris Dyer, Yoav Goldberg, Graham Neubig

Neural Nets and Language

Language

Discrete, structured (graphs, trees)

Big challenge: writing code that translates between the
{discrete-structured, continuous} regimes

Neural-Nets

Continuous: poor native support for
structure

Why not do it yourself?

- Hard to compare with existing models
- Obscures difference between model and optimization
- Debugging has to be custom-built
- Hard to tweak model

Outline

- Computation graphs (general)
- Neural Nets in PyTorch
- Full example

Computation Graphs

Expression

\vec{x}


graph:



Computation Graphs

Expression

\vec{x}^\top




- Edge: function argument / data dependency
- A node with an incoming edge is a function $F \equiv f(u)$ edge's tail node
- A node computes its value and the value of its derivative w.r.t each argument (edge) times a derivative $\frac{\partial f}{\partial u}$

Computation Graphs

Expression

$$\vec{x}^\top A$$

graph:




Functions can be nullary, unary, binary, . . . n-ary. Often they are unary or binary.

Computation Graphs

Expression

$$\vec{x}^\top A x$$

graph:





Computation graphs are (usually) directed and acyclic

Computation Graphs

Expression

$$\vec{x}^\top A x$$

graph:




$$\frac{\partial f(\mathbf{x}, \mathbf{A})}{\partial \mathbf{x}} = (\mathbf{A}^\top + \mathbf{A})\mathbf{x}$$
$$\frac{\partial f(\mathbf{x}, \mathbf{A})}{\partial \mathbf{A}} = \mathbf{x}\mathbf{x}^\top$$

Computation Graphs

Expression

$$\vec{x}^\top A x + b \cdot \vec{x} + c$$

graph:




Computation Graphs

Expression

$$y = \vec{x}^\top A x + b \cdot \vec{x} + c$$

graph:




Variable names label nodes


Algorithms

- Graph construction
- Forward propagation
 - Loop over nodes in topological order
 - Compute the value of the node given its inputs
 - Given my inputs, make a prediction (i.e. “error” vs. “target output”)
- Backward propagation
 - Loop over the nodes in reverse topological order, starting with goal node
 - Compute derivatives of final goal node value wrt each edge's tail node
 - How does the output change with small change to inputs?


Forward Propagation




Forward Propagation




Forward Propagation




Forward Propagation




Forward Propagation




Forward Propagation



Forward Propagation



Forward Propagation



Constructing Graphs

Static declaration

- Define architecture, run data through
- PROS: Optimization, hardware support
- CONS: Structured data ugly, graph language

Theano, Tensorflow

Dynamic declaration

- Graph implicit with data
- PROS: Native language, interleave construction/evaluation
- CONS: Slower, computation can be wasted

Chainer, Dynet, PyTorch

Constructing Graphs

Static declaration

- Define architecture, run data through
- PROS: Optimization, hardware support
- CONS: Structured data ugly, graph language


Theano, Tensorflow

Dynamic declaration


- Graph implicit with data
- PROS: Native language, interleave construction/evaluation
- CONS: Slower, computation can be wasted

Chainer, Dynet, PyTorch


Words



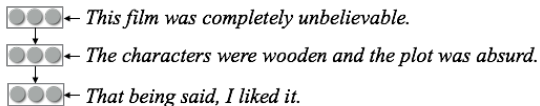
Sentences



Phrases



Documents



Language is Hierarchical

Dynamic Hierarchy in Language

- Language is hierarchical
 - Graph should reflect this reality
 - Traditional flow-control best for processing
- Combinatorial algorithms (e.g., dynamic programming)
- Exploit independencies to compute over a large space of operations tractably


PyTorch

- Torch: Facebook's deep learning framework
- Nice, but written in Lua (C backend)
- Optimized to run computations on GPU
- Mature, industry-supported framework

Why GPU?



Why GPU?





Frameworks

Natural Language Processing: Jordan
Boyd-Graber

University of Maryland

BACKPROP IN PYTORCH

Simple Model

```
import torch
import torch.nn as nn

class LogisticRegression(nn.Module):
    def __init__(self, input_size, num_classes):
        super(LogisticRegression, self).__init__()
        self.linear = nn.Linear(input_size, num_classes)

    def forward(self, x):
        out = self.linear(x)
        return out
```


Simple Model

```
>>> model = LogisticRegression(5, 2)
>>> model.parameters
<bound method Module.parameters of LogisticRegression(
  (linear): Linear(in_features=5, out_features=2, bias=True)
)>
>>> model.linear.weight
Parameter containing:
tensor([[ 0.0650,  0.0221,  0.1673, -0.1365, -0.1233],
        [-0.1289,  0.2455,  0.3255,  0.0409, -0.1908]], requires_grad=True)
>>> model.linear.bias
Parameter containing:
tensor([-0.2208,  0.2562], requires_grad=True)
```

Where did these numbers come from?

```
class Bilinear(Module):  
    r"""Applies a bilinear transformation to the incoming o  
    :math: 'y = x_1 A x_2 + b '  
    """  
  
    def reset_parameters(self):  
        stdv = 1. / math.sqrt(self.weight.size(1))  
        self.weight.data.uniform_(-stdv, stdv)  
        if self.bias is not None:  
            self.bias.data.uniform_(-stdv, stdv)
```

Where did these numbers come from?

```
class Bilinear(Module):  
    r"""Applies a bilinear transformation to the incoming o  
    :math: 'y = x_1 A x_2 + b '  
    """  
  
    def reset_parameters(self):  
        stdv = 1. / math.sqrt(self.weight.size(1))  
        self.weight.data.uniform_(-stdv, stdv)  
        if self.bias is not None:  
            self.bias.data.uniform_(-stdv, stdv)
```

Beauty and peril of working with something like PyTorch!

Computation Graph and Expressions

- Create basic expressions.
- Combine them using operations.
- Expressions represent symbolic computations.
- Actual computation:

```
.value()  
.npvalue()           #numpy value  
.scalar_value()  
.cuda()              # move to GPU  
.forward()           # compute expression
```

Running Computation Forward

```
>>> x = torch.Tensor(1, 5)
>>> x
tensor([[ 0.0000, -0.0000,  0.0000, -0.0000,  0.0000]])
>>> x = x*0 + 1
>>> x
tensor([[1., 1., 1., 1., 1.]])
>>> model.forward(x)
tensor([[ -0.2263,  0.5485]]), grad_fn=<ThAddmmBackward>)
```

Modules allow computation graph

- Each module must implement forward function
- If forward function just uses built-in modules, autograd works
- If not, you'll need to implement backward function (i.e., backprop)

Modules allow computation graph

- Each module must implement forward function
- If forward function just uses built-in modules, autograd works
- If not, you'll need to implement backward function (i.e., backprop)
 - input: as many Tensors as outputs of module (gradient w.r.t. that output)
 - output: as many Tensors as inputs of module (gradient w.r.t. its corresponding input)
 - If inputs do not need gradient (static) you can return None

Trainers and Backprop

- Initialize a Optimizer with a given model's parameter
- Get output for an example / minibatch
- Compute loss and backpropagate
- Take step of Optimizer
- Repeat ...

Trainers and Backprop

```
optimizer = torch.optim.SGD(model.parameters(),  
                             lr=learning_rate)  
  
# Training the Model  
for epoch in range(num_epochs):  
    for i, (Variable(doc), Variable(label)) in \  
        enumerate(train_loader):  
        optimizer.zero_grad()  
        prediction = model(doc)  
        loss = nn.CrossEntropyLoss(prediction, label)  
        loss.backward()  
        optimizer.step()
```

Options for Optimizers

Adadelat

Adagrad

Adam

LBFGS

SGD

Closure (LBFGS), learning rate, etc.

Key Points

- Create computation graph for each example.
- Graph is built by composing expressions.
- Functions that take expressions and return expressions define graph components.

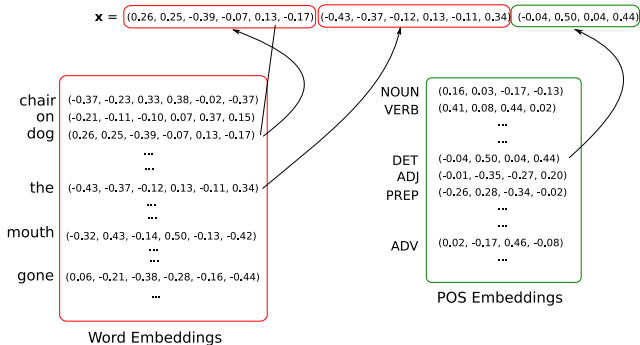
Word Embeddings and Lookup Parameters

- In NLP, it is very common to use feature embeddings
- Each feature is represented as a d -dim vector
- These are then summed or concatenated to form an input vector
- The embeddings can be pre-trained
- But they are usually trained (fine-tuned) with the model

"feature embeddings"

$w=\text{dog}$ $pw=\text{the}$ $pt=\text{NOUN}$ $pt=\text{DET}$ $w=\text{dog}\&pt=\text{DET}$ $w=\text{dog}\&pw=\text{the}$ $w=\text{chair}\&pt=\text{DET}$

$\mathbf{x} = (0, \dots, 0, 1, 0, \dots, 0, 1, 0, \dots, 0, 1, 0, \dots, 0, 1, 0, 0, 1, 0, \dots, 0, 0, 0, \dots, 0)$



```
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim

torch.manual_seed(1)
word_to_ix = {"hello": 0, "world": 1}
embeds = nn.Embedding(2, 5)  # 2 words in vocab, 5 dim embeddings
lookup_tensor = torch.tensor([word_to_ix["hello"]],
                              dtype=torch.long)
hello_embed = embeds(lookup_tensor)
```



Frameworks

Natural Language Processing: Jordan
Boyd-Graber

University of Maryland

EXAMPLE IMPLEMENTATION: DAN

Deep Unordered Composition Rivals Syntactic Methods for Text Classification

Mohit Iyyer,¹ Varun Manjunatha,¹ Jordan Boyd-Graber,² Hal Daumé III¹

¹University of Maryland, Department of Computer Science and UMIACS

²University of Colorado, Department of Computer Science

{miyyer, varunm, hal}@umiacs.umd.edu, Jordan.Boyd.Graber@colorado.edu

Implementing a non-trivial example . . .

Deep Averaging Network

w_1, \dots, w_N



$z_0 = \text{CBOW}(w_1, \dots, w_N)$

$z_1 = g(W_1 z_0 + b_1)$

$z_2 = g(W_2 z_1 + b_2)$

$\hat{y} = \text{softmax}(z_2)$

- Works about as well as more complicated models
- Strong baseline
- Key idea: Continuous Bag of Words

$$\text{CBOW}(w_1, \dots, w_N) = \sum_i E[w_i] \quad (1)$$

- Actual non-linearity doesn't matter, we'll use tanh
- Let's implement in PyTorch

Deep Averaging Network

$$\begin{aligned} &w_1, \dots, w_N \\ &\quad \downarrow \\ &z_0 = \text{CBOW}(w_1, \dots, w_N) \\ &z_1 = g(z_0) \\ &z_2 = g(z_1) \\ &\hat{y} = \text{softmax}(z_2) \end{aligned}$$

Initialization

```
def __init__(self, n_classes, vocab_size, emb_dim=300,
              n_hidden_units=300):
    super(DanModel, self).__init__()
    self.n_classes = n_classes
    self.vocab_size = vocab_size
    self.emb_dim = emb_dim
    self.n_hidden_units = n_hidden_units
    self.embeddings = nn.Embedding(self.vocab_size,
                                    self.emb_dim)
    self.classifier = nn.Sequential(
        nn.Linear(self.n_hidden_units,
                   self.n_hidden_units),
        nn.ReLU(),
        nn.Linear(self.n_hidden_units,
                   self.n_classes))
    self._softmax = nn.Softmax()
```

Deep Averaging Network

$$\begin{aligned} &w_1, \dots, w_N \\ &\quad \downarrow \\ &z_0 = \text{CBOW}(w_1, \dots, w_N) \\ &z_1 = g(z_0) \\ &z_2 = g(z_1) \\ &\hat{y} = \text{softmax}(z_2) \end{aligned}$$

Forward

```
def forward(self, batch, probs=False):
    text = batch['text']['tokens']
    length = batch['length']
    text_embed = self._word_embeddings(text)
    # Take the mean embedding. Since padding results
    # in zeros its safe to sum and divide by length
    encoded = text_embed.sum(1)
    encoded /= lengths.view(text_embed.size(0), -1)

    # Compute the network score predictions
    logits = self.classifier(encoded)
    if probs:
        return self._softmax(logits)
    else:
        return logits
```

Deep Averaging Network

w_1, \dots, w_N



$z_0 = \text{CBOW}(w_1, \dots, w_N)$

$z_1 = g(z_0)$

$z_2 = g(z_1)$

$\hat{y} = \text{softmax}(z_2)$

Training

```
def _run_epoch(self, batch_iter, train=True):
    self._model.train()
    for batch in batch_iter:
        model.zero_grad()
        out = model(batches)
        batch_loss = criterion(out,
                                batch['label'])
        batch_loss.backward()
        self.optimizer.step()
```

Summary

- Computation Graph
- Expressions (\approx nodes in the graph)
- Parameters, LookupParameters
- Model (a collection of parameters)
- Optimizers
- Create a graph for each example, compute loss, backdrop, update



Multilayer Networks

Natural Language Processing: Jordan
Boyd-Graber

University of Maryland

HANDS-ON DEMO

Data and Model

Data

x_1	x_2	y
1.00	1.00	0.00
1.00	0.00	1.00
0.00	0.00	0.00
0.00	1.00	1.00

First Layer

$$w^{(1)} = \begin{bmatrix} 1.00 & 1.00 \\ 1.00 & 1.00 \end{bmatrix} \quad (1)$$

$$b^{(1)} = [-1.00 \quad 0.00] \quad (2)$$

Second Layer

$$w^{(2)} = [-2.00 \quad 1.00] \quad (3)$$

$$b^{(2)} = 0.00 \quad (4)$$

Using ReLU as non-linearity

Prediction for $x_0 = (1.00, 1.00)$

Prediction for $x_0 = (1.00, 1.00)$

■ Hidden Computation

$$a_{0,0}^{(1)} = f(w_{0,0}^{(1)} \cdot 1.00 + w_{0,1}^{(1)} \cdot 1.00 + b_0) \quad (5)$$

$$= f(1.00 \cdot 1.00 + 1.00 \cdot 1.00 + -1.00) \quad (6)$$

Prediction for $x_0 = (1.00, 1.00)$

■ Hidden Computation

$$a_{0,0}^{(1)} = f(w_{0,0}^{(1)} \cdot 1.00 + w_{0,1}^{(1)} \cdot 1.00 + b_0) \quad (5)$$

$$= f(1.00 \cdot 1.00 + 1.00 \cdot 1.00 + -1.00) \quad (6)$$

$$a_{0,1}^{(1)} = f(w_{1,0}^{(1)} \cdot 1.00 + w_{1,1}^{(1)} \cdot 1.00 + b_1) \quad (7)$$

$$= f(1.00 \cdot 1.00 + 1.00 \cdot 1.00 + 0.00) \quad (8)$$

Prediction for $x_0 = (1.00, 1.00)$

■ Hidden Computation

$$a_{0,0}^{(1)} = f(w_{0,0}^{(1)} \cdot 1.00 + w_{0,1}^{(1)} \cdot 1.00 + b_0) \quad (5)$$

$$= f(1.00 \cdot 1.00 + 1.00 \cdot 1.00 + -1.00) \quad (6)$$

$$a_{0,1}^{(1)} = f(w_{1,0}^{(1)} \cdot 1.00 + w_{1,1}^{(1)} \cdot 1.00 + b_1) \quad (7)$$

$$= f(1.00 \cdot 1.00 + 1.00 \cdot 1.00 + 0.00) \quad (8)$$

■ Hidden Layer: [1 2]

■ Output Answer

$$a_{0,0}^{(3)} = f(w_{0,0}^{(2)} \cdot 1.00 + w_{0,1}^{(2)} \cdot 2.00 + b_0) \quad (9)$$

$$= f(-2.00 \cdot 1.00 + 1.00 \cdot 2.00 + 0.00) \quad (10)$$

■ Prediction: 0.00, Error: 0.00

Prediction for $x_0 = (1.00, 1.00)$

■ Hidden Computation

$$a_{0,0}^{(1)} = f(w_{0,0}^{(1)} \cdot 1.00 + w_{0,1}^{(1)} \cdot 1.00 + b_0) \quad (5)$$

$$= f(1.00 \cdot 1.00 + 1.00 \cdot 1.00 + -1.00) \quad (6)$$

$$a_{0,1}^{(1)} = f(w_{1,0}^{(1)} \cdot 1.00 + w_{1,1}^{(1)} \cdot 1.00 + b_1) \quad (7)$$

$$= f(1.00 \cdot 1.00 + 1.00 \cdot 1.00 + 0.00) \quad (8)$$

■ Hidden Layer: [1 2]

■ Output Answer

$$a_{0,0}^{(3)} = f(w_{0,0}^{(2)} \cdot 1.00 + w_{0,1}^{(2)} \cdot 2.00 + b_0) \quad (9)$$

$$= f(-2.00 \cdot 1.00 + 1.00 \cdot 2.00 + 0.00) \quad (10)$$

Prediction for $x_0 = (1.00, 1.00)$

- Hidden Computation

$$a_{0,0}^{(1)} = f(w_{0,0}^{(1)} \cdot 1.00 + w_{0,1}^{(1)} \cdot 1.00 + b_0) \quad (5)$$

$$= f(1.00 \cdot 1.00 + 1.00 \cdot 1.00 + -1.00) \quad (6)$$

$$a_{0,1}^{(1)} = f(w_{1,0}^{(1)} \cdot 1.00 + w_{1,1}^{(1)} \cdot 1.00 + b_1) \quad (7)$$

$$= f(1.00 \cdot 1.00 + 1.00 \cdot 1.00 + 0.00) \quad (8)$$

- Hidden Layer: [1 2]

- Output Answer

$$a_{0,0}^{(3)} = f(w_{0,0}^{(2)} \cdot 1.00 + w_{0,1}^{(2)} \cdot 2.00 + b_0) \quad (9)$$

$$= f(-2.00 \cdot 1.00 + 1.00 \cdot 2.00 + 0.00) \quad (10)$$

Prediction for $x_1 = (1.00, 0.00)$

Prediction for $x_1 = (1.00, 0.00)$

■ Hidden Computation

$$a_{1,0}^{(1)} = f(w_{0,0}^{(1)} \cdot 1.00 + w_{0,1}^{(1)} \cdot 0.00 + b_0) \quad (11)$$

$$= f(1.00 \cdot 1.00 + 1.00 \cdot 0.00 + -1.00) \quad (12)$$

Prediction for $x_1 = (1.00, 0.00)$

■ Hidden Computation

$$a_{1,0}^{(1)} = f(w_{0,0}^{(1)} \cdot 1.00 + w_{0,1}^{(1)} \cdot 0.00 + b_0) \quad (11)$$

$$= f(1.00 \cdot 1.00 + 1.00 \cdot 0.00 + -1.00) \quad (12)$$

$$a_{1,1}^{(1)} = f(w_{1,0}^{(1)} \cdot 1.00 + w_{1,1}^{(1)} \cdot 0.00 + b_1) \quad (13)$$

$$= f(1.00 \cdot 1.00 + 1.00 \cdot 0.00 + 0.00) \quad (14)$$

Prediction for $x_1 = (1.00, 0.00)$

■ Hidden Computation

$$a_{1,0}^{(1)} = f(w_{0,0}^{(1)} \cdot 1.00 + w_{0,1}^{(1)} \cdot 0.00 + b_0) \quad (11)$$

$$= f(1.00 \cdot 1.00 + 1.00 \cdot 0.00 + -1.00) \quad (12)$$

$$a_{1,1}^{(1)} = f(w_{1,0}^{(1)} \cdot 1.00 + w_{1,1}^{(1)} \cdot 0.00 + b_1) \quad (13)$$

$$= f(1.00 \cdot 1.00 + 1.00 \cdot 0.00 + 0.00) \quad (14)$$

■ Hidden Layer: [0 1]

■ Output Answer

$$a_{1,0}^{(3)} = f(w_{0,0}^{(2)} \cdot 0.00 + w_{0,1}^{(2)} \cdot 1.00 + b_0) \quad (15)$$

$$= f(-2.00 \cdot 0.00 + 1.00 \cdot 1.00 + 0.00) \quad (16)$$

■ Prediction: 1.00, Error: 0.00

Prediction for $x_1 = (1.00, 0.00)$

- Hidden Computation

$$a_{1,0}^{(1)} = f(w_{0,0}^{(1)} \cdot 1.00 + w_{0,1}^{(1)} \cdot 0.00 + b_0) \quad (11)$$

$$= f(1.00 \cdot 1.00 + 1.00 \cdot 0.00 + -1.00) \quad (12)$$

$$a_{1,1}^{(1)} = f(w_{1,0}^{(1)} \cdot 1.00 + w_{1,1}^{(1)} \cdot 0.00 + b_1) \quad (13)$$

$$= f(1.00 \cdot 1.00 + 1.00 \cdot 0.00 + 0.00) \quad (14)$$

- Hidden Layer: [0 1]

- Output Answer

$$a_{1,0}^{(3)} = f(w_{0,0}^{(2)} \cdot 0.00 + w_{0,1}^{(2)} \cdot 1.00 + b_0) \quad (15)$$

$$= f(-2.00 \cdot 0.00 + 1.00 \cdot 1.00 + 0.00) \quad (16)$$

Prediction for $x_1 = (1.00, 0.00)$

- Hidden Computation

$$a_{1,0}^{(1)} = f(w_{0,0}^{(1)} \cdot 1.00 + w_{0,1}^{(1)} \cdot 0.00 + b_0) \quad (11)$$

$$= f(1.00 \cdot 1.00 + 1.00 \cdot 0.00 + -1.00) \quad (12)$$

$$a_{1,1}^{(1)} = f(w_{1,0}^{(1)} \cdot 1.00 + w_{1,1}^{(1)} \cdot 0.00 + b_1) \quad (13)$$

$$= f(1.00 \cdot 1.00 + 1.00 \cdot 0.00 + 0.00) \quad (14)$$

- Hidden Layer: [0 1]

- Output Answer

$$a_{1,0}^{(3)} = f(w_{0,0}^{(2)} \cdot 0.00 + w_{0,1}^{(2)} \cdot 1.00 + b_0) \quad (15)$$

$$= f(-2.00 \cdot 0.00 + 1.00 \cdot 1.00 + 0.00) \quad (16)$$

Prediction for $x_2 = (0.00, 0.00)$

Prediction for $x_2 = (0.00, 0.00)$

■ Hidden Computation

$$a_{2,0}^{(1)} = f(w_{0,0}^{(1)} \cdot 0.00 + w_{0,1}^{(1)} \cdot 0.00 + b_0) \quad (17)$$

$$= f(1.00 \cdot 0.00 + 1.00 \cdot 0.00 + -1.00) \quad (18)$$

Prediction for $x_2 = (0.00, 0.00)$

■ Hidden Computation

$$a_{2,0}^{(1)} = f(w_{0,0}^{(1)} \cdot 0.00 + w_{0,1}^{(1)} \cdot 0.00 + b_0) \quad (17)$$

$$= f(1.00 \cdot 0.00 + 1.00 \cdot 0.00 + -1.00) \quad (18)$$

$$a_{2,1}^{(1)} = f(w_{1,0}^{(1)} \cdot 0.00 + w_{1,1}^{(1)} \cdot 0.00 + b_1) \quad (19)$$

$$= f(1.00 \cdot 0.00 + 1.00 \cdot 0.00 + 0.00) \quad (20)$$

Prediction for $x_2 = (0.00, 0.00)$

■ Hidden Computation

$$a_{2,0}^{(1)} = f(w_{0,0}^{(1)} \cdot 0.00 + w_{0,1}^{(1)} \cdot 0.00 + b_0) \quad (17)$$

$$= f(1.00 \cdot 0.00 + 1.00 \cdot 0.00 + -1.00) \quad (18)$$

$$a_{2,1}^{(1)} = f(w_{1,0}^{(1)} \cdot 0.00 + w_{1,1}^{(1)} \cdot 0.00 + b_1) \quad (19)$$

$$= f(1.00 \cdot 0.00 + 1.00 \cdot 0.00 + 0.00) \quad (20)$$

■ Hidden Layer: [0. 0.]

■ Output Answer

$$a_{2,0}^{(3)} = f(w_{0,0}^{(2)} \cdot 0.00 + w_{0,1}^{(2)} \cdot 0.00 + b_0) \quad (21)$$

$$= f(-2.00 \cdot 0.00 + 1.00 \cdot 0.00 + 0.00) \quad (22)$$

■ Prediction: 0.00, Error: 0.00

Prediction for $x_2 = (0.00, 0.00)$

- Hidden Computation

$$a_{2,0}^{(1)} = f(w_{0,0}^{(1)} \cdot 0.00 + w_{0,1}^{(1)} \cdot 0.00 + b_0) \quad (17)$$

$$= f(1.00 \cdot 0.00 + 1.00 \cdot 0.00 + -1.00) \quad (18)$$

$$a_{2,1}^{(1)} = f(w_{1,0}^{(1)} \cdot 0.00 + w_{1,1}^{(1)} \cdot 0.00 + b_1) \quad (19)$$

$$= f(1.00 \cdot 0.00 + 1.00 \cdot 0.00 + 0.00) \quad (20)$$

- Hidden Layer: [0. 0.]

- Output Answer

$$a_{2,0}^{(3)} = f(w_{0,0}^{(2)} \cdot 0.00 + w_{0,1}^{(2)} \cdot 0.00 + b_0) \quad (21)$$

$$= f(-2.00 \cdot 0.00 + 1.00 \cdot 0.00 + 0.00) \quad (22)$$

Prediction for $x_2 = (0.00, 0.00)$

- Hidden Computation

$$a_{2,0}^{(1)} = f(w_{0,0}^{(1)} \cdot 0.00 + w_{0,1}^{(1)} \cdot 0.00 + b_0) \quad (17)$$

$$= f(1.00 \cdot 0.00 + 1.00 \cdot 0.00 + -1.00) \quad (18)$$

$$a_{2,1}^{(1)} = f(w_{1,0}^{(1)} \cdot 0.00 + w_{1,1}^{(1)} \cdot 0.00 + b_1) \quad (19)$$

$$= f(1.00 \cdot 0.00 + 1.00 \cdot 0.00 + 0.00) \quad (20)$$

- Hidden Layer: [0. 0.]

- Output Answer

$$a_{2,0}^{(3)} = f(w_{0,0}^{(2)} \cdot 0.00 + w_{0,1}^{(2)} \cdot 0.00 + b_0) \quad (21)$$

$$= f(-2.00 \cdot 0.00 + 1.00 \cdot 0.00 + 0.00) \quad (22)$$

Prediction for $x_3 = (0.00, 1.00)$

Prediction for $x_3 = (0.00, 1.00)$

■ Hidden Computation

$$a_{3,0}^{(1)} = f(w_{0,0}^{(1)} \cdot 0.00 + w_{0,1}^{(1)} \cdot 1.00 + b_0) \quad (23)$$

$$= f(1.00 \cdot 0.00 + 1.00 \cdot 1.00 + -1.00) \quad (24)$$

Prediction for $x_3 = (0.00, 1.00)$

■ Hidden Computation

$$a_{3,0}^{(1)} = f(w_{0,0}^{(1)} \cdot 0.00 + w_{0,1}^{(1)} \cdot 1.00 + b_0) \quad (23)$$

$$= f(1.00 \cdot 0.00 + 1.00 \cdot 1.00 + -1.00) \quad (24)$$

$$a_{3,1}^{(1)} = f(w_{1,0}^{(1)} \cdot 0.00 + w_{1,1}^{(1)} \cdot 1.00 + b_1) \quad (25)$$

$$= f(1.00 \cdot 0.00 + 1.00 \cdot 1.00 + 0.00) \quad (26)$$

Prediction for $x_3 = (0.00, 1.00)$

- Hidden Computation

$$a_{3,0}^{(1)} = f(w_{0,0}^{(1)} \cdot 0.00 + w_{0,1}^{(1)} \cdot 1.00 + b_0) \quad (23)$$

$$= f(1.00 \cdot 0.00 + 1.00 \cdot 1.00 + -1.00) \quad (24)$$

$$a_{3,1}^{(1)} = f(w_{1,0}^{(1)} \cdot 0.00 + w_{1,1}^{(1)} \cdot 1.00 + b_1) \quad (25)$$

$$= f(1.00 \cdot 0.00 + 1.00 \cdot 1.00 + 0.00) \quad (26)$$

- Hidden Layer: [0 1]

- Output Answer

$$a_{3,0}^{(3)} = f(w_{0,0}^{(2)} \cdot 0.00 + w_{0,1}^{(2)} \cdot 1.00 + b_0) \quad (27)$$

$$= f(-2.00 \cdot 0.00 + 1.00 \cdot 1.00 + 0.00) \quad (28)$$

- Prediction: 1.00, Error: 0.00

Prediction for $x_3 = (0.00, 1.00)$

■ Hidden Computation

$$a_{3,0}^{(1)} = f(w_{0,0}^{(1)} \cdot 0.00 + w_{0,1}^{(1)} \cdot 1.00 + b_0) \quad (23)$$

$$= f(1.00 \cdot 0.00 + 1.00 \cdot 1.00 + -1.00) \quad (24)$$

$$a_{3,1}^{(1)} = f(w_{1,0}^{(1)} \cdot 0.00 + w_{1,1}^{(1)} \cdot 1.00 + b_1) \quad (25)$$

$$= f(1.00 \cdot 0.00 + 1.00 \cdot 1.00 + 0.00) \quad (26)$$

■ Hidden Layer: [0 1]

■ Output Answer

$$a_{3,0}^{(3)} = f(w_{0,0}^{(2)} \cdot 0.00 + w_{0,1}^{(2)} \cdot 1.00 + b_0) \quad (27)$$

$$= f(-2.00 \cdot 0.00 + 1.00 \cdot 1.00 + 0.00) \quad (28)$$

Prediction for $x_3 = (0.00, 1.00)$

■ Hidden Computation

$$a_{3,0}^{(1)} = f(w_{0,0}^{(1)} \cdot 0.00 + w_{0,1}^{(1)} \cdot 1.00 + b_0) \quad (23)$$

$$= f(1.00 \cdot 0.00 + 1.00 \cdot 1.00 + -1.00) \quad (24)$$

$$a_{3,1}^{(1)} = f(w_{1,0}^{(1)} \cdot 0.00 + w_{1,1}^{(1)} \cdot 1.00 + b_1) \quad (25)$$


$$= f(1.00 \cdot 0.00 + 1.00 \cdot 1.00 + 0.00) \quad (26)$$


■ Hidden Layer: [0 1]

■ Output Answer

$$a_{3,0}^{(3)} = f(w_{0,0}^{(2)} \cdot 0.00 + w_{0,1}^{(2)} \cdot 1.00 + b_0) \quad (27)$$

$$= f(-2.00 \cdot 0.00 + 1.00 \cdot 1.00 + 0.00) \quad (28)$$





Next Time

- Representing Words
- Updating representations
- Comparing with contextual information